



deeplearning.ai

Optimization Algorithms

Mini-batch
gradient descent

Batch vs. mini-batch gradient descent

Vectorization allows you to efficiently compute on m examples.

Mini-batch gradient descent




deeplearning.ai

Optimization Algorithms


Understanding
mini-batch
gradient descent

Training with mini batch gradient descent


Batch gradient descent



Mini-batch gradient descent



Choosing your mini-batch size



Choosing your mini-batch size



deeplearning.ai

Optimization Algorithms

Exponentially weighted averages

Temperature in London

$$\theta_1 = 40^\circ\text{F}$$

$$\theta_2 = 49^\circ\text{F}$$


$$\theta_3 = 45^\circ\text{F}$$

\vdots


$$\theta_{180} = 60^\circ\text{F}$$

$$\theta_{181} = 56^\circ\text{F}$$

\vdots



Exponentially weighted averages






deeplearning.ai

Optimization Algorithms

Understanding
exponentially
weighted averages

Exponentially weighted averages

$$v_t = \beta v_{t-1} + (1 - \beta)\theta_t$$



Exponentially weighted averages

$$v_t = \beta v_{t-1} + (1 - \beta)\theta_t$$

$$v_{100} = 0.9v_{99} + 0.1\theta_{100}$$

$$v_{99} = 0.9v_{98} + 0.1\theta_{99}$$

$$v_{98} = 0.9v_{97} + 0.1\theta_{98}$$

...

Implementing exponentially weighted averages

$$v_0 = 0$$

$$v_1 = \beta v_0 + (1 - \beta) \theta_1$$

$$v_2 = \beta v_1 + (1 - \beta) \theta_2$$

$$v_3 = \beta v_2 + (1 - \beta) \theta_3$$

...




deeplearning.ai

Optimization Algorithms

Bias correction
in exponentially
weighted average

Bias correction



$$v_t = \beta v_{t-1} + (1 - \beta) \theta_t$$




deeplearning.ai

Optimization Algorithms

Gradient descent with momentum

Gradient descent example



deeplearning.ai

Implementation details

On iteration t :

Compute dW, db on the current mini-batch

$$v_{dW} = \beta v_{dW} + (1 - \beta) dW$$

$$v_{db} = \beta v_{db} + (1 - \beta) db$$

$$W = W - \alpha v_{dW}, \quad b = b - \alpha v_{db}$$

Hyperparameters: α, β $\beta = 0.9$




deeplearning.ai

Optimization Algorithms

RMSprop

RMSprop





deeplearning.ai

Optimization Algorithms

Adam optimization algorithm

Adam optimization algorithm

```
yhat = np.array([.9, 0.2, 0.1, .4, .9])
```

Hyperparameters choice:



Adam Coates




deeplearning.ai

Optimization Algorithms

Learning rate decay

Learning rate decay



Learning rate decay

Other learning rate decay methods





deeplearning.ai


Optimization Algorithms

The problem of local optima

Local optima in neural networks



Problem of plateaus



- Unlikely to get stuck in a bad local optima
- Plateaus can make learning slow